

nba_regression_analysis

February 21, 2024

0.1 Factors impacting an NBA players 2023-2024 guaranteed contract value analysis notebook.

What data are we exploring? I am interested in sports analytics, and my favorite sport is basketball, so I used player performance stats from the last 5 NBA seasons. I wrote custom functions that leverage player names and IDs to pull stats from the NBA_API. After executing the functions and merging the returned dataframes, we have a dataset that contains 476 rows and 21 columns. The columns represent standard NBA performance metrics such as: - total points - total steals - total assists - etc.

This notebook aims to inform players and their management firms if or how player performance impacts contract value and focus on visualization and analysis. *The data collection and pre-processing was completed in a separate notebook to maintain clarity as you follow along.*

0.1.1 1. Import libraries required for the analysis

Below is a list of libraries used to analyze and visualize the data.

```
[ ]: import math
import numpy as np
import pandas as pd
import plotly
import plotly.express as px
import plotly.io as pio
pio.renderers.default = "notebook+pdf"
from sklearn.preprocessing import StandardScaler
import statsmodels.api as sm
import sys
import warnings

pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)
warnings.simplefilter(action='ignore', category=FutureWarning)
```

C:\Users\toobr\AppData\Local\Temp\ipykernel_20272\3289393469.py:3:

DeprecationWarning:

Pyarrow will become a required dependency of pandas in the next major release of pandas (pandas 3.0),

(to allow more performant data types, such as the Arrow string type, and better

interoperability with other libraries)
but was not found to be installed on your system.
If this would cause problems for you,
please provide us feedback at <https://github.com/pandas-dev/pandas/issues/54466>

```
import pandas as pd
```

0.1.2 1. Load Data For Exploratory Analysis

```
[ ]: #read in pre-processed file and view sample
preprocessed_data = pd.read_csv('../data/nba_regression_analysis_df.csv')
preprocessed_data.head()
```

```
[ ]: Current_Contract First_Name Last_Name PLAYER_ID GP GS MIN FGM \
0 51915615.0 stephen curry 201939 257 257 8774.0 2417
1 47649433.0 kevin durant 201142 184 181 6548.0 1868
2 47607350.0 nikola jokic 203999 368 368 12126.0 3323
3 47607350.0 joel embiid 203954 300 300 9825.0 2817
4 47607350.0 lebron james 2544 278 277 9795.0 2872

FGA FG3M FG3A FTM FTA OREB DREB AST STL BLK TOV PF PTS
0 5144 1261 3049 1183 1287 151 1265 1513 311 78 806 543 7278
1 3442 386 949 1196 1315 76 1204 1016 141 231 623 381 5318
2 5903 409 1199 1597 1930 972 3220 2953 484 257 1224 1007 8652
3 5561 353 1040 2659 3206 673 2796 1089 284 467 984 894 8646
4 5661 645 1862 1225 1717 280 1972 2212 321 185 1000 491 7614
```

The purpose of exploratory analysis is to understand the properties of the data, discover patterns, and determine if there are additional data quality issues that need to be addressed.

```
[ ]: #confirm datatypes
preprocessed_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 419 entries, 0 to 418
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Current_Contract      419 non-null   float64
1   First_Name            419 non-null   object
2   Last_Name             419 non-null   object
3   PLAYER_ID             419 non-null   int64
4   GP                    419 non-null   int64
5   GS                    419 non-null   int64
6   MIN                   419 non-null   float64
7   FGM                   419 non-null   int64
8   FGA                   419 non-null   int64
9   FG3M                  419 non-null   int64
10  FG3A                  419 non-null   int64
```

```

11  FTM          419 non-null    int64
12  FTA          419 non-null    int64
13  OREB         419 non-null    int64
14  DREB         419 non-null    int64
15  AST          419 non-null    int64
16  STL          419 non-null    int64
17  BLK          419 non-null    int64
18  TOV          419 non-null    int64
19  PF           419 non-null    int64
20  PTS          419 non-null    int64

```

dtypes: float64(2), int64(17), object(2)

memory usage: 68.9+ KB

```
[ ]: #convert Player ID to string
preprocessed_data['PLAYER_ID'] = preprocessed_data['PLAYER_ID'].astype(str)
```

```
[ ]: #review the summary statistics
preprocessed_data.describe().round()
```

```
[ ]:
```

	Current_Contract	GP	GS	MIN	FGM	FGA	FG3M	\
count	419.0	419.0	419.0	419.0	419.0	419.0	419.0	
mean	11052897.0	198.0	112.0	4929.0	876.0	1859.0	254.0	
std	11670322.0	94.0	105.0	3235.0	735.0	1539.0	253.0	
min	92857.0	10.0	0.0	62.0	7.0	19.0	0.0	
25%	2368860.0	123.0	20.0	2127.0	319.0	683.0	46.0	
50%	6175000.0	206.0	72.0	4634.0	684.0	1428.0	163.0	
75%	15552322.0	273.0	199.0	7444.0	1236.0	2788.0	426.0	
max	51915615.0	377.0	368.0	12485.0	3428.0	6741.0	1261.0	

	FG3A	FTM	FTA	OREB	DREB	AST	STL	BLK	TOV	\
count	419.0	419.0	419.0	419.0	419.0	419.0	419.0	419.0	419.0	
mean	700.0	383.0	490.0	209.0	701.0	536.0	155.0	100.0	283.0	
std	664.0	428.0	530.0	209.0	588.0	572.0	118.0	108.0	255.0	
min	0.0	0.0	0.0	1.0	9.0	4.0	0.0	0.0	2.0	
25%	150.0	98.0	136.0	72.0	262.0	135.0	56.0	32.0	98.0	
50%	486.0	246.0	322.0	145.0	571.0	341.0	136.0	63.0	213.0	
75%	1154.0	524.0	646.0	266.0	958.0	726.0	220.0	135.0	390.0	
max	3049.0	2659.0	3434.0	1374.0	3448.0	3285.0	559.0	748.0	1461.0	

	PF	PTS
count	419.0	419.0
mean	396.0	2390.0
std	249.0	2048.0
min	8.0	21.0
25%	184.0	830.0
50%	375.0	1808.0
75%	572.0	3424.0

```
max    1087.0  9529.0
```

I'm curious about the following and will use visualizations to dig deeper: - How many games have the top earners played? - How does the number of games played relate to contract values? - How does the number of games started relate to contract values?

```
[ ]: #create a dataframe of top 5 highest earners
top_earners = preprocessed_data.nlargest(10, 'Current_Contract').
    ↪sort_values('Current_Contract', ascending= False)

# Create a horizontal bar chart
fig = px.bar(top_earners, x='GP', y='First_Name', orientation='h',
             title='How Many Games Have the Highest Earners Played?',
             labels={'GP': 'Number of Games Played Over 5 Seasons',
                    'First_Name': 'Player Name'})

# Update layout
fig.update_layout(width=800, height=400,
                  xaxis_title='Number of Games Played Over 5 Seasons',
                  yaxis_title='Player Name')

# Show the plot
fig.show()
```

With the exception of Kevin Durant and Kawhi Leonard the top earners have all played at least 3 seasons over the last five years. Next I will explore the relationship between the number of games played and contract values.

```
[ ]: #print the correlation metric
print(preprocessed_data[['GP', 'Current_Contract']].corr())

# Create scatter plot
fig = px.scatter(preprocessed_data, x='GP', y='Current_Contract',
                 title='Relationship between Current Contract and Games Played',
                 labels={'GP': 'Games Played', 'Current_Contract': 'Current_
    ↪Contract'})

# Update layout
fig.update_layout(width=800, height=600)

# Show the plot
fig.show()
```

	GP	Current_Contract
GP	1.000000	0.566553
Current_Contract	0.566553	1.000000

The scatterplot and correlation value tells me that there isn't strong relationship between games played and contract values. In fact the scatter plot shows the a majority of players making less than

\$10 million this year. This makes sense in the context of the data because NBA teams usually have 12-man rosters but will only play 7-8 guys. Let's see if Games started has a strong relationship to contract values.

```
[ ]: #print the correlation metric
print(preprocessed_data[['GS', 'Current_Contract']].corr())

# Create scatter plot
fig = px.scatter(preprocessed_data, x='GS', y='Current_Contract',
                 title='Relationship between Current Contract and Games_
↳Started',
                 labels={'GS': 'Games Started', 'Current_Contract': 'Current_
↳Contract'})

# Update layout
fig.update_layout(width=800, height=600)

# Show the plot
fig.show()
```

	GS	Current_Contract
GS	1.000000	0.781069
Current_Contract	0.781069	1.000000

Okay, so we can see a stronger relationship between games started and contract values which makes sense in the context of the data because players who start are generally more valuable and demand higher salaries. Let's investigate the distribution of games played in the cell below

```
[ ]: num_bins = math.ceil(1 + math.log2(len(preprocessed_data)))

# Create histogram plot
fig = px.histogram(preprocessed_data, x='GP', nbins=num_bins,
                  title='Distribution of Games Played',
                  labels={'GP': 'Games Played', 'count': 'Frequency'})

# Update layout
fig.update_layout(width=500, height=600)

# Show the plot
fig.show()
```

Referring back to the descriptive statistics there's a sizable difference between Min and Max Games Played ('GP'). Additionally, the chart in the cell above shows that games played is skewed. I will create per game stats so that players who play less games can be compared to players who play more games.

```
[ ]: #iterate over the column calculating per game stats
per_game_cols = preprocessed_data.columns[-16:]
```

```
for col in per_game_cols:
    new_col_name = f'{col}_Per_Game'
    preprocessed_data[new_col_name] = round(preprocessed_data[col] /
                                             preprocessed_data['GP'],2)
```

```
[ ]: preprocessed_data.head().round()
```

```
[ ]:
Current_Contract First_Name Last_Name PLAYER_ID GP GS MIN FGM \
0 51915615.0 stephen curry 201939 257 257 8774.0 2417
1 47649433.0 kevin durant 201142 184 181 6548.0 1868
2 47607350.0 nikola jokic 203999 368 368 12126.0 3323
3 47607350.0 joel embiid 203954 300 300 9825.0 2817
4 47607350.0 lebron james 2544 278 277 9795.0 2872

FGA FG3M FG3A FTM FTA OREB DREB AST STL BLK TOV PF PTS \
0 5144 1261 3049 1183 1287 151 1265 1513 311 78 806 543 7278
1 3442 386 949 1196 1315 76 1204 1016 141 231 623 381 5318
2 5903 409 1199 1597 1930 972 3220 2953 484 257 1224 1007 8652
3 5561 353 1040 2659 3206 673 2796 1089 284 467 984 894 8646
4 5661 645 1862 1225 1717 280 1972 2212 321 185 1000 491 7614

GS_Per_Game MIN_Per_Game FGM_Per_Game FGA_Per_Game FG3M_Per_Game \
0 1.0 34.0 9.0 20.0 5.0
1 1.0 36.0 10.0 19.0 2.0
2 1.0 33.0 9.0 16.0 1.0
3 1.0 33.0 9.0 19.0 1.0
4 1.0 35.0 10.0 20.0 2.0

FG3A_Per_Game FTM_Per_Game FTA_Per_Game OREB_Per_Game DREB_Per_Game \
0 12.0 5.0 5.0 1.0 5.0
1 5.0 6.0 7.0 0.0 7.0
2 3.0 4.0 5.0 3.0 9.0
3 3.0 9.0 11.0 2.0 9.0
4 7.0 4.0 6.0 1.0 7.0

AST_Per_Game STL_Per_Game BLK_Per_Game TOV_Per_Game PF_Per_Game \
0 6.0 1.0 0.0 3.0 2.0
1 6.0 1.0 1.0 3.0 2.0
2 8.0 1.0 1.0 3.0 3.0
3 4.0 1.0 2.0 3.0 3.0
4 8.0 1.0 1.0 4.0 2.0

PTS_Per_Game
0 28.0
1 29.0
2 24.0
3 29.0
```

4 27.0

Now that we have per game stats, I will drop the aggregated total columns.

```
[ ]: totals_col = ['GP', 'GS', 'MIN', 'FGM', 'FGA', 'FG3M', 'FG3A', 'FTM', 'FTA',  
                  'OREB', 'DREB', 'AST', 'STL', 'BLK', 'TOV', 'PF', 'PTS']
```

```
regression_df = preprocessed_data.drop(columns= totals_col)
```

```
[ ]: regression_df.head()
```

```
[ ]:      Current_Contract  First_Name  Last_Name  PLAYER_ID  GS_Per_Game  MIN_Per_Game  \  
0          51915615.0      stephen      curry      201939          1.00          34.14  
1          47649433.0        kevin      durant      201142          0.98          35.59  
2          47607350.0      nikola      jokic      203999          1.00          32.95  
3          47607350.0        joel      embiid      203954          1.00          32.75  
4          47607350.0      lebron      james       2544          1.00          35.23  
  
      FGM_Per_Game  FGA_Per_Game  FG3M_Per_Game  FG3A_Per_Game  FTM_Per_Game  \  
0           9.40          20.02           4.91          11.86           4.60  
1          10.15          18.71           2.10           5.16           6.50  
2           9.03          16.04           1.11           3.26           4.34  
3           9.39          18.54           1.18           3.47           8.86  
4          10.33          20.36           2.32           6.70           4.41  
  
      FTA_Per_Game  OREB_Per_Game  DREB_Per_Game  AST_Per_Game  STL_Per_Game  \  
0           5.01           0.59           4.92           5.89           1.21  
1           7.15           0.41           6.54           5.52           0.77  
2           5.24           2.64           8.75           8.02           1.32  
3          10.69           2.24           9.32           3.63           0.95  
4           6.18           1.01           7.09           7.96           1.15  
  
      BLK_Per_Game  TOV_Per_Game  PF_Per_Game  PTS_Per_Game  
0           0.30           3.14           2.11          28.32  
1           1.26           3.39           2.07          28.90  
2           0.70           3.33           2.74          23.51  
3           1.56           3.28           2.98          28.82  
4           0.67           3.60           1.77          27.39
```

```
[ ]: regression_df['PTS_Per_Game'].describe().round(2)
```

```
[ ]: count      419.00  
     mean       10.52  
     std        6.06  
     min        1.24  
     25%        5.94  
     50%        9.01  
     75%       14.07
```

```
max          29.23
Name: PTS_Per_Game, dtype: float64
```

In related studies, assists, points, and turnovers have impacted contract values, and while my dataset is different from the related studies I am curious to understand how those metric relate to contract values.

```
[ ]: #create a new column called 'AST_Rate' that categorizes 'AST_per_game'
conditions = [
    regression_df['AST_Per_Game'].round(2) >= 4.20,
    (regression_df['AST_Per_Game'].round(2) < 4.20) &
    (regression_df['AST_Per_Game'].round(2) >= 0.44),
    regression_df['AST_Per_Game'].round(2) < 0.44
]
values = ['high', 'medium', 'low']

# Create the 'AST_Rate' column based on the conditions
regression_df['AST_Rate'] = np.select(conditions, values, default=np.nan)

# Create boxplot
fig = px.box(regression_df, x='AST_Rate', y='Current_Contract',
             title='A Comparison of Assist Per Game',
             color='AST_Rate',
             color_discrete_map={'high': '#064096', 'medium': '#81a6c8',
                                'low': '#adc89e'},
             labels={'AST_Rate': 'Assist Rate',
                    'Current_Contract': 'Contract Value'})

# Update layout
fig.update_layout(width=400,
                  height=400,
                  xaxis=dict(showgrid=False),
                  yaxis=dict(showgrid=False))

# Show the plot
fig.show()
```

In the cell above we can see that players with high assist rates (>4.2 per game) tend to earn more money. This sample does include the entire population so the medium and low rate are representative.

```
[ ]: #create a new column called 'TO_Rate' that categorizes 'TOV_per_game'
conditions = [
    regression_df['TOV_Per_Game'].round(2) >= 2.07,
    (regression_df['TOV_Per_Game'].round(2) < 2.07) &
    (regression_df['TOV_Per_Game'].round(2) >= 0.47),
    regression_df['TOV_Per_Game'].round(2) < 0.47
]
```



```

values = ['high', 'medium', 'low']

# Create the 'AST_Rate' column based on the conditions
regression_df['TO_Rate'] = np.select(conditions, values, default=np.nan)

# Create boxplot
fig = px.scatter(regression_df, x='TOV_Per_Game', y='Current_Contract',
                 title='Comparison of Distributions by Turnover Rate',
                 color='TO_Rate',
                 color_discrete_map={'high': '#064096', 'medium': '#81a6c8',
                                     'low': '#adc89e'},
                 labels={'TO_Rate': 'Turnover Rate',
                        'Current_Contract': 'Contract Value'})

# Update layout
fig.update_layout(width=500,
                  height=400,
                  xaxis=dict(showgrid=False),
                  yaxis=dict(showgrid=False))

# Show the plot
fig.show()

# print the correlation metric
print(regression_df[['TOV_Per_Game', 'Current_Contract']].corr())

```

	TOV_Per_Game	Current_Contract
TOV_Per_Game	1.000000	0.713049
Current_Contract	0.713049	1.000000

The cell above shows a moderately strong positive correlation between turnovers and contract values. This is a case of correlation does not equal causation, because turnovers actually have a negative impact on the game. When a player turns the ball over his team loses an opportunity to score points. In the context of this dataset, players with higher contract values also have higher turnover rates because they usually have the ball more.

```

[ ]: #create a new column called 'Score_Rate' that categorizes 'PTS_per_game'
conditions = [
    regression_df['PTS_Per_Game'].round(2) >= 16.58,
    (regression_df['PTS_Per_Game'].round(2) < 16.58) &
    (regression_df['PTS_Per_Game'].round(2) >= 4.46),
    regression_df['PTS_Per_Game'].round(2) < 4.46
]
values = ['high', 'medium', 'low']

# Create the 'AST_Rate' column based on the conditions
regression_df['Score_Rate'] = np.select(conditions, values, default=np.nan)

```

```

# Create boxplot
fig = px.scatter(regression_df, x='PTS_Per_Game', y='Current_Contract',
                 title='Comparison of Distributions by Scoring Rate',
                 color='Score_Rate',
                 color_discrete_map={'high': '#064096', 'medium': '#81a6c8',
                                     'low': '#adc89e'},
                 labels={'Score_Rate': 'Scoring Rate',
                         'Current_Contract': 'Current Contract'})

# Update layout
fig.update_layout(width=500,
                  height=400,
                  xaxis=dict(showgrid=False),
                  yaxis=dict(showgrid=False))

# Show the plot
fig.show()

# print the correlation metric
print(regression_df[['PTS_Per_Game', 'Current_Contract']].corr())

```

	PTS_Per_Game	Current_Contract
PTS_Per_Game	1.000000	0.841468
Current_Contract	0.841468	1.000000

As expected there is a strong positive correlation between the scoring rate and contract values, in the context of this dataset higher scorers are more valuable because the team with the most points will win the game.

Are there any columns that lack variance necessary to make predictions?

```
[ ]: print(regression_df.var(numeric_only= True))
```

Current_Contract	1.361964e+14
GS_Per_Game	1.243046e-01
MIN_Per_Game	6.168737e+01
FGM_Per_Game	4.611111e+00
FGA_Per_Game	2.027354e+01
FG3M_Per_Game	7.049936e-01
FG3A_Per_Game	4.716682e+00
FTM_Per_Game	1.984625e+00
FTA_Per_Game	2.993166e+00
OREB_Per_Game	5.369617e-01
DREB_Per_Game	2.865559e+00
AST_Per_Game	3.528240e+00
STL_Per_Game	1.260525e-01
BLK_Per_Game	1.543743e-01
TOV_Per_Game	6.398220e-01
PF_Per_Game	3.965185e-01

```
PTS_Per_Game      3.671359e+01
dtype: float64
```

0.1.3 2. Check Regression Assumptions

Before we conduct regression analysis we need to check that the data set meets the following assumptions:

- Linearity: The relationship between independent and dependent variables is linear
- Independence: The observations are independent of each other

```
[ ]: #use a SPLOM chart to visually check for linear relationship
fig = px.scatter_matrix(regression_df,
                        dimensions=['Current_Contract', 'GS_Per_Game',
                                   'MIN_Per_Game', 'FGM_Per_Game',
                                   'FGA_Per_Game', 'FG3M_Per_Game',
                                   'FG3A_Per_Game', 'FTM_Per_Game',
                                   'FTA_Per_Game', 'OREB_Per_Game',
                                   'DREB_Per_Game', 'AST_Per_Game',
                                   'STL_Per_Game', 'BLK_Per_Game',
                                   'TOV_Per_Game', 'PF_Per_Game',
                                   'PTS_Per_Game'])

# Update layout
fig.update_layout(
    title='SPLOM',
    width=1200,
    height=1200,
    font=dict(size=6)
)
fig.show()
```

I focused on the top row of the SPLOM chart to visually inspect if a linear relationship exists between Current_Contract and the other variables in the dataset. To confirm I will also print the Pearson correlation coefficient.

```
[ ]: #print Pearson's R to validate visual
regression_df.corr(numeric_only=True)['Current_Contract']
```

```
[ ]: Current_Contract      1.000000
     GS_Per_Game           0.705001
     MIN_Per_Game          0.723979
     FGM_Per_Game          0.833019
     FGA_Per_Game          0.799279
     FG3M_Per_Game         0.508811
     FG3A_Per_Game         0.509879
     FTM_Per_Game          0.776162
     FTA_Per_Game          0.758861
     OREB_Per_Game         0.190251
     DREB_Per_Game         0.597306
```

```

AST_Per_Game      0.618480
STL_Per_Game      0.530077
BLK_Per_Game      0.271679
TOV_Per_Game      0.713049
PF_Per_Game       0.444709
PTS_Per_Game      0.841468
Name: Current_Contract, dtype: float64

```

While some relationships are stronger than others, each variable has evidence of linearity. Next I will check for correlation.

```

[ ]: #use correlation matrix to visualize correlations
correlation_matrix = regression_df.corr(numeric_only= True)

```

```

[ ]: #create heatmap
fig = px.imshow(correlation_matrix,
                 x=correlation_matrix.index,
                 y=correlation_matrix.columns,
                 color_continuous_scale='temps',
                 zmin= -1,
                 zmax = 1,
                 text_auto= True)

# Update layout
fig.update_layout(
    title='Correlation Heatmap',
    width=800,
    height=800
)

# Show plot
fig.show()

```

The following features have strong correlations to other features in the dataset and will be removed:

- GS_Per_Game - MIN_Per_Game - FGA_Per_Game - FG3A_Per_Game - FTA_Per_Game
- PTS_Per_Game - FTM_Per_Game - TOV_Per_Game

```

[ ]: #drop correlated columns
correlated_cols = ['GS_Per_Game', 'MIN_Per_Game', 'FGA_Per_Game', 'FG3A_Per_Game',
                  'FTA_Per_Game', 'PTS_Per_Game', 'FTM_Per_Game', 'TOV_Per_Game']

regression_df.drop(columns= correlated_cols, inplace= True)

```

Re-check correlations

```

[ ]: #use correlation matrix to visualize correlations
correlation_matrix = regression_df.corr(numeric_only= True)

#create heatmap

```

```
fig = px.imshow(correlation_matrix,
                 x=correlation_matrix.index,
                 y=correlation_matrix.columns,
                 color_continuous_scale='temps',
                 zmin= -1,
                 zmax = 1,
                 text_auto= True)

# Update layout
fig.update_layout(
    title='Correlation Heatmap',
    width=800,
    height=800
)

# Show plot
fig.show()
```

```
[ ]: #save a copy of the final dataframe used for OLS
regression_df.to_csv('../data/regression_df.csv', index = False)
```

0.1.4 3. Regression Analysis

Now, I will use the regression_df to understand how performance metrics impact contract values.

```
[ ]: #view the columns
regression_df.columns
```

```
[ ]: Index(['Current_Contract', 'First_Name', 'Last_Name', 'PLAYER_ID',
           'FGM_Per_Game', 'FG3M_Per_Game', 'OREB_Per_Game', 'DREB_Per_Game',
           'AST_Per_Game', 'STL_Per_Game', 'BLK_Per_Game', 'PF_Per_Game',
           'AST_Rate', 'TO_Rate', 'Score_Rate'],
          dtype='object')
```

```
[ ]: #Declare X and y variables
X = regression_df[['FGM_Per_Game', 'FG3M_Per_Game', 'OREB_Per_Game',
                  'DREB_Per_Game', 'AST_Per_Game', 'STL_Per_Game',
                  'BLK_Per_Game', 'PF_Per_Game']]
y = regression_df['Current_Contract']

#Scale to variables to account for the difference in magnitude
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_scaled_df = pd.DataFrame(X_scaled, columns=X.columns)
df_scaled = pd.concat([X_scaled_df, y], axis=1)

#Fit the OLS regression model using the formula API
formula = 'Current_Contract ~ ' + ' + '.join(df_scaled.columns[:-1])
```

```
ols_model = sm.formula.ols(formula=formula, data=df_scaled).fit()

#Print the model summary
print(ols_model.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          Current_Contract      R-squared:                0.711
Model:                  OLS                   Adj. R-squared:           0.706
Method:                 Least Squares         F-statistic:              126.2
Date:                   Wed, 21 Feb 2024       Prob (F-statistic):       1.46e-105
Time:                   09:20:27              Log-Likelihood:           -7152.1
No. Observations:       419                   AIC:                     1.432e+04
Df Residuals:           410                   BIC:                     1.436e+04
Df Model:                8
Covariance Type:        nonrobust
=====
```

```
=====
=
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-
Intercept      1.105e+07    3.09e+05    35.725    0.000    1.04e+07
1.17e+07
FGM_Per_Game   7.981e+06    6.75e+05    11.820    0.000    6.65e+06
9.31e+06
FG3M_Per_Game  8.301e+05    5.29e+05     1.568    0.118   -2.11e+05
1.87e+06
OREB_Per_Game -6.528e+05    6.28e+05    -1.039    0.299   -1.89e+06
5.82e+05
DREB_Per_Game  9.273e+05    6.41e+05     1.446    0.149   -3.34e+05
2.19e+06
AST_Per_Game   1.163e+06    5.72e+05     2.033    0.043    3.85e+04
2.29e+06
STL_Per_Game   2.098e+05    4.89e+05     0.429    0.668   -7.52e+05
1.17e+06
BLK_Per_Game   1.603e+06    4.62e+05     3.472    0.001    6.96e+05
2.51e+06
PF_Per_Game    -1.072e+06    4.64e+05    -2.310    0.021   -1.98e+06
-1.6e+05
=====
```

```
Omnibus:                22.489    Durbin-Watson:           1.286
Prob(Omnibus):           0.000    Jarque-Bera (JB):        39.013
Skew:                    -0.352    Prob(JB):                 3.38e-09
Kurtosis:                 4.319    Cond. No.                  5.71
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

From the output above we can see 4 features have significance (P-value >.05). Let's fit a model using the significant features only.

```
[ ]: #Declare X and y variables
X_final = regression_df[['FGM_Per_Game', 'AST_Per_Game', 'BLK_Per_Game',
                        'PF_Per_Game']]
y_final = regression_df['Current_Contract']

#Scale to variables to account for the difference in magnitude
scaler = StandardScaler()
X_scaled_final = scaler.fit_transform(X_final)
X_scaled_df_final = pd.DataFrame(X_scaled_final, columns=X_final.columns)
df_scaled_final = pd.concat([X_scaled_df_final, y_final], axis=1)

#Fit the OLS regression model using the formula API
formula_final = 'Current_Contract ~ ' + ' + '.join(df_scaled_final.columns[:-1])
ols_model_final = sm.formula.ols(formula=formula_final,
                                data=df_scaled_final).fit()

#Print the model summary
print(ols_model_final.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          Current_Contract      R-squared:                0.705
Model:                  OLS                  Adj. R-squared:            0.702
Method:                 Least Squares         F-statistic:              247.5
Date:                   Wed, 21 Feb 2024      Prob (F-statistic):       2.40e-108
Time:                   09:20:30              Log-Likelihood:           -7156.4
No. Observations:       419                  AIC:                     1.432e+04
Df Residuals:           414                  BIC:                     1.434e+04
Df Model:               4
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.105e+07	3.11e+05	35.530	0.000	1.04e+07	1.17e+07
FGM_Per_Game	8.953e+06	4.94e+05	18.113	0.000	7.98e+06	9.92e+06
AST_Per_Game	1.385e+06	4.59e+05	3.019	0.003	4.83e+05	2.29e+06
BLK_Per_Game	1.264e+06	3.87e+05	3.271	0.001	5.05e+05	2.02e+06
PF_Per_Game	-9.875e+05	4.26e+05	-2.316	0.021	-1.83e+06	-1.49e+05

```
=====
Omnibus:                 20.094      Durbin-Watson:              1.286
Prob(Omnibus):            0.000      Jarque-Bera (JB):           37.867
Skew:                     -0.281     Prob(JB):                   5.99e-09
Kurtosis:                 4.361      Cond. No.                   3.03
=====
```

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Interpret Results: - FGM_Per_Game: Increasing average FGs by 1 adds nearly 9 million dollars to contract value. - AST_Per_Game: Increasing average assists by 1 adds 1.3 million dollars to contract value. - BLK_Per_Game: Increasing average blocks by 1 adds 1.2 million dollars to contract value. - PF_Per_Game: Increasing average fouls by 1 subtracts 9.8 million dollars from contract value.